

# Intelligent Agents


## Chapter 2

# Outline

- Agents and environments
- Rationality
- Task environment:  
    *PEAS:*
  - *Performance measure*
  - *Environment*
  - *Actuators*
  - *Sensors*
- Environment types
- Agent types


# Agents and Environments

- An *agent* is anything that can be viewed as perceiving its *environment* through *sensors* and acting in that environment through *actuators*.



# Agents and Environments

- An *agent* is anything that can be viewed as perceiving its *environment* through *sensors* and acting in that environment through *actuators*.




- Agents* include humans, robots, softbots, thermostats, etc.
- The *agent function* maps from percept histories to actions:

$$f : \mathcal{P}^* \rightarrow \mathcal{A}$$

- The *agent program* runs on a physical *architecture* to give  $f$

## Vacuum-cleaner world



Percepts: location and contents, e.g.,  $[A, \textit{Dirty}]$

Actions: *Left*, *Right*, *Suck*, *NoOp*

## A vacuum-cleaner agent

Agent function:

Percept sequence	Action
[A, <i>Clean</i> ]	<i>Right</i>
[A, <i>Dirty</i> ]	<i>Suck</i>
[B, <i>Clean</i> ]	<i>Left</i>
[B, <i>Dirty</i> ]	<i>Suck</i>
[A, <i>Clean</i> ], [A, <i>Clean</i> ]	<i>Right</i>
[A, <i>Clean</i> ], [A, <i>Dirty</i> ]	<i>Suck</i>
...	...

Note: This says *how* the agent should function.

- It says nothing about how this should be implemented.

# A vacuum-cleaner agent

Agent program:

Function **Reflex-Vacuum-Agent**([location,status]) returns an action  
    if status = Dirty then return Suck  
    else if location = A then return Right  
    else if location = B then return Left

Ask:

- What is the *right* function for implementing a specification?
- Can it be implemented in a small agent program?

# Rationality

Informally a *rational* agent is one that does the “right thing”.



# Rationality

Informally a *rational* agent is one that does the “right thing”.

- *How well* an agent does is given by a *performance measure*.

# Rationality

Informally a *rational* agent is one that does the “right thing”.

- *How well* an agent does is given by a *performance measure*.
- A fixed *performance measure* evaluates a sequence of environment states
- Examples:
  - one point per square cleaned up in time  $T$ ?
  - one point per clean square per time step, minus one per move?
  - penalize for  $> k$  dirty squares?

# Rationality

Informally a *rational* agent is one that does the “right thing”.

- How well an agent does is given by a *performance measure*.
- A fixed *performance measure* evaluates a sequence of environment states
- Examples:
  - one point per square cleaned up in time  $T$ ?
  - one point per clean square per time step, minus one per move?
  - penalize for  $> k$  dirty squares?
- A *rational agent* selects an action which maximizes the **expected** value of the performance measure **given the percept sequence to date** and its own knowledge.
- The action selection may range from being hardwired (e.g. in an insect or reflexive agent) to involving substantial reasoning.

# Rationality

Notes:

- Rational  $\neq$  omniscient
  - percepts may not supply all the relevant information

# Rationality

## Notes:

- Rational  $\neq$  omniscient
  - percepts may not supply all the relevant information
- Rational  $\neq$  clairvoyant
  - action outcomes may not be as expected

# Rationality

## Notes:

- Rational  $\neq$  omniscient
  - percepts may not supply all the relevant information
- Rational  $\neq$  clairvoyant
  - action outcomes may not be as expected
- Hence, rational  $\neq$  successful
- Full, general rationality requires exploration, learning, autonomy

# The Task Environment

- To design a rational agent, we must specify the *task environment*
- The task environment has the following components:
  - Performance measure
  - Environment
  - Actuators
  - Sensors
- Acronym: PEAS

# PEAS

Consider, e.g., the task of designing an automated taxi:

**Performance measure:** safety, destination, profits, legality, comfort,  
...

**Environment:** streets/freeways, traffic, pedestrians, weather, ...

**Actuators:** steering, accelerator, brake, horn, speaker, ...

**Sensors:** video, accelerometers, gauges, engine sensors,  
keyboard, GPS, ...



# Internet shopping agent

Performance measure: ??

Environment: ??

Actuators: ??

Sensors: ??

# Internet shopping agent

Performance measure: price, quality, appropriateness, efficiency

Environment: ??

Actuators: ??

Sensors: ??

# Internet shopping agent

Performance measure: price, quality, appropriateness, efficiency

Environment: current and future WWW sites, vendors, shippers

Actuators: ??

Sensors: ??

# Internet shopping agent

**Performance measure:** price, quality, appropriateness, efficiency

**Environment:** current and future WWW sites, vendors, shippers

**Actuators:** display to user, follow URL, fill in form

**Sensors:** ??

# Internet shopping agent

**Performance measure:** price, quality, appropriateness, efficiency

**Environment:** current and future WWW sites, vendors, shippers

**Actuators:** display to user, follow URL, fill in form

**Sensors:** HTML pages (text, graphics, scripts)

# Environment Types

- *Fully observable vs. partially observable*
  - If the agent has access to full state of the environment or not

# Environment Types

- *Fully observable vs. partially observable*
  - If the agent has access to full state of the environment or not
- *Deterministic vs. stochastic*
  - Deterministic: Next state is completely determined by the agent's actions. (Or the set of agents in a multiagent env.)

# Environment Types

- *Fully observable vs. partially observable*
  - If the agent has access to full state of the environment or not
- *Deterministic vs. stochastic*
  - Deterministic: Next state is completely determined by the agent's actions. (Or the set of agents in a multiagent env.)

👉 *Uncertain*: not fully observable or not deterministic



# Environment Types

- *Fully observable vs. partially observable*
  - If the agent has access to full state of the environment or not
- *Deterministic vs. stochastic*
  - Deterministic: Next state is completely determined by the agent's actions. (Or the set of agents in a multiagent env.)
- 👉 *Uncertain*: not fully observable or not deterministic
- *Episodic vs. sequential*
  - Episodic: Agent's experience is divided into independent episodes (e.g. classification)

# Environment Types

- *Fully observable vs. partially observable*
  - If the agent has access to full state of the environment or not
- *Deterministic vs. stochastic*
  - Deterministic: Next state is completely determined by the agent's actions. (Or the set of agents in a multiagent env.)
- ☞ *Uncertain*: not fully observable or not deterministic
- *Episodic vs. sequential*
  - Episodic: Agent's experience is divided into independent episodes (e.g. classification)
- *Static vs. dynamic*
  - Dynamic: Environment may change while agent is deliberating.

# Environment Types

- *Fully observable vs. partially observable*
  - If the agent has access to full state of the environment or not
- *Deterministic vs. stochastic*
  - Deterministic: Next state is completely determined by the agent's actions. (Or the set of agents in a multiagent env.)

👉 *Uncertain*: not fully observable or not deterministic

- *Episodic vs. sequential*
  - Episodic: Agent's experience is divided into independent episodes (e.g. classification)
- *Static vs. dynamic*
  - Dynamic: Environment may change while agent is deliberating.
- *Discrete vs. continuous*

# Environment Types

- *Fully observable vs. partially observable*
  - If the agent has access to full state of the environment or not
- *Deterministic vs. stochastic*
  - Deterministic: Next state is completely determined by the agent's actions. (Or the set of agents in a multiagent env.)
- 👉 *Uncertain: not fully observable or not deterministic*
- *Episodic vs. sequential*
  - Episodic: Agent's experience is divided into independent episodes (e.g. classification)
- *Static vs. dynamic*
  - Dynamic: Environment may change while agent is deliberating.
- *Discrete vs. continuous*
- *Single-agent vs. multiagent*

## Environment types

	Crossword	Backgammon	Internet shopping	Taxi
Observable				
Deterministic				
Episodic				
Static				
Discrete				
Single-agent				

## Environment types

	Crossword	Backgammon	Internet shopping	Taxi
Observable	Yes	Yes	No	No
Deterministic				
Episodic				
Static				
Discrete				
Single-agent				

## Environment types

	Crossword	Backgammon	Internet shopping	Taxi
Observable	Yes	Yes	No	No
Deterministic	Yes	No	Partly	No
Episodic				
Static				
Discrete				
Single-agent				

## Environment types

	Crossword	Backgammon	Internet shopping	Taxi
Observable	Yes	Yes	No	No
Deterministic	Yes	No	Partly	No
Episodic	No	No	No	No
Static				
Discrete				
Single-agent				



## Environment types

	Crossword	Backgammon	Internet shopping	Taxi
Observable	Yes	Yes	No	No
Deterministic	Yes	No	Partly	No
Episodic	No	No	No	No
Static	Yes	Yes	Semi	No
Discrete				
Single-agent				

## Environment types

	Crossword	Backgammon	Internet shopping	Taxi
Observable	Yes	Yes	No	No
Deterministic	Yes	No	Partly	No
Episodic	No	No	No	No
Static	Yes	Yes	Semi	No
Discrete	Yes	Yes	Yes	No
Single-agent				

## Environment types

	Crossword	Backgammon	Internet shopping	Taxi
Observable	Yes	Yes	No	No
Deterministic	Yes	No	Partly	No
Episodic	No	No	No	No
Static	Yes	Yes	Semi	No
Discrete	Yes	Yes	Yes	No
Single-agent	Yes	No	Yes	No

## Environment types

	Crossword	Backgammon	Internet shopping	Taxi
Observable	Yes	Yes	No	No
Deterministic	Yes	No	Partly	No
Episodic	No	No	No	No
Static	Yes	Yes	Semi	No
Discrete	Yes	Yes	Yes	No
Single-agent	Yes	No	Yes	No

👉 *The environment type largely determines the agent design*

- The real world is:

## Environment types

	Crossword	Backgammon	Internet shopping	Taxi
Observable	Yes	Yes	No	No
Deterministic	Yes	No	Partly	No
Episodic	No	No	No	No
Static	Yes	Yes	Semi	No
Discrete	Yes	Yes	Yes	No
Single-agent	Yes	No	Yes	No

👉 *The environment type largely determines the agent design*

- The real world is:
  - partially observable,  
stochastic,  
sequential,  
dynamic,  
continuous, and  
multi-agent

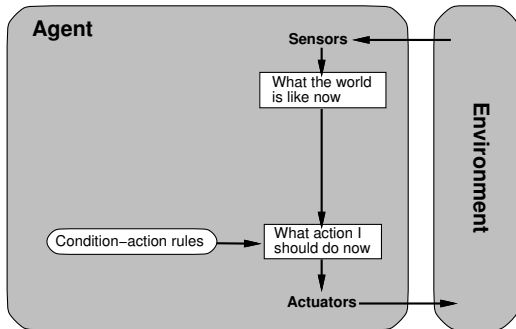
## Agent types

There are four basic types in order of increasing generality:

- simple reflex agents
- reflex agents with state
- goal-based agents
- utility-based agents

All these can have a learning component added

## Simple reflex agents



- Action is selected according to the current percept
- No knowledge of percept history.

## A simple reflex agent algorithm

Function **Simple-Reflex-Agent**(percept) **returns** an action  
**persistent:** rules a set of condition-action rules

state  $\leftarrow$  Interpret-Input(percept)

rule  $\leftarrow$  Rule-Match(state,rules)

action  $\leftarrow$  rule.Action

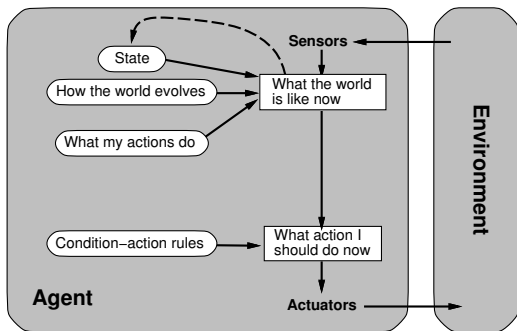
**return** action



## Example

Function **Reflex-Vacuum-Agent**([location,status]) returns an action  
if status = Dirty then return Suck  
    else if location = A then return Right  
    else if location = B then return Left

## Reflex agents with state



- Also called a “model-based reflex agent”
- Agent keeps track of what it knows about the world.
- Useful for partial observability

## A simple reflex agent algorithm

Function **Reflex-Agent-With-State**(**percept**) **returns** an action

- persistent:** **state**: the agent's conception of the world state
- model**: The transition model – how the next state depends on the present state and action
- rules**: a set of condition-action rules
- action**: the most recent action (initially none)

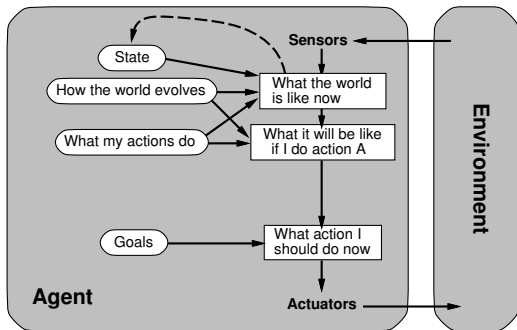
**state**  $\leftarrow$  Update-State(**state**,**action**,**percept**,**model**)

**rule**  $\leftarrow$  Rule-Match(**state**,**rules**)

**action**  $\leftarrow$  **rule**.Action

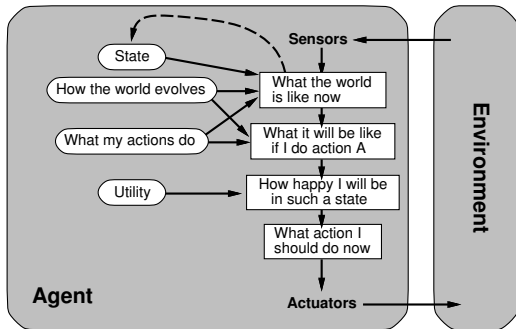
**return** **action**

## Goal-based agents



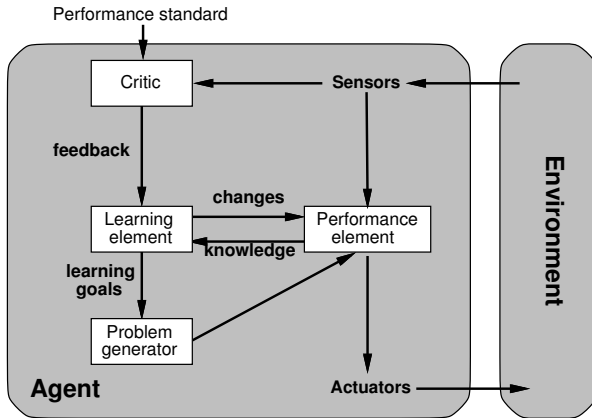
- Agent's actions are determined in part by its *goals*.
- Example: Classical planning.

## Utility-based agents



- In addition to goals, use a notion of how “good” an action sequence is.
  - E.g.: Taxi to airport should be safe, efficient, etc.

# Learning agents



## Summary

- *Agents* interact with *environments* through *actuators* and *sensors*
- The *agent function* describes what the agent does in all circumstances
- The *performance measure* evaluates the environment sequence
- A *rational* agent maximizes expected performance
- *Agent programs* implement agent functions
- *PEAS* descriptions define task environments
- Environments are categorized along several dimensions:  
*observable? deterministic? episodic? static? discrete?*  
*single-agent?*
- Several basic agent architectures exist:  
*reflex, reflex with state, goal-based, utility-based*