

Graphs

CMPT 225

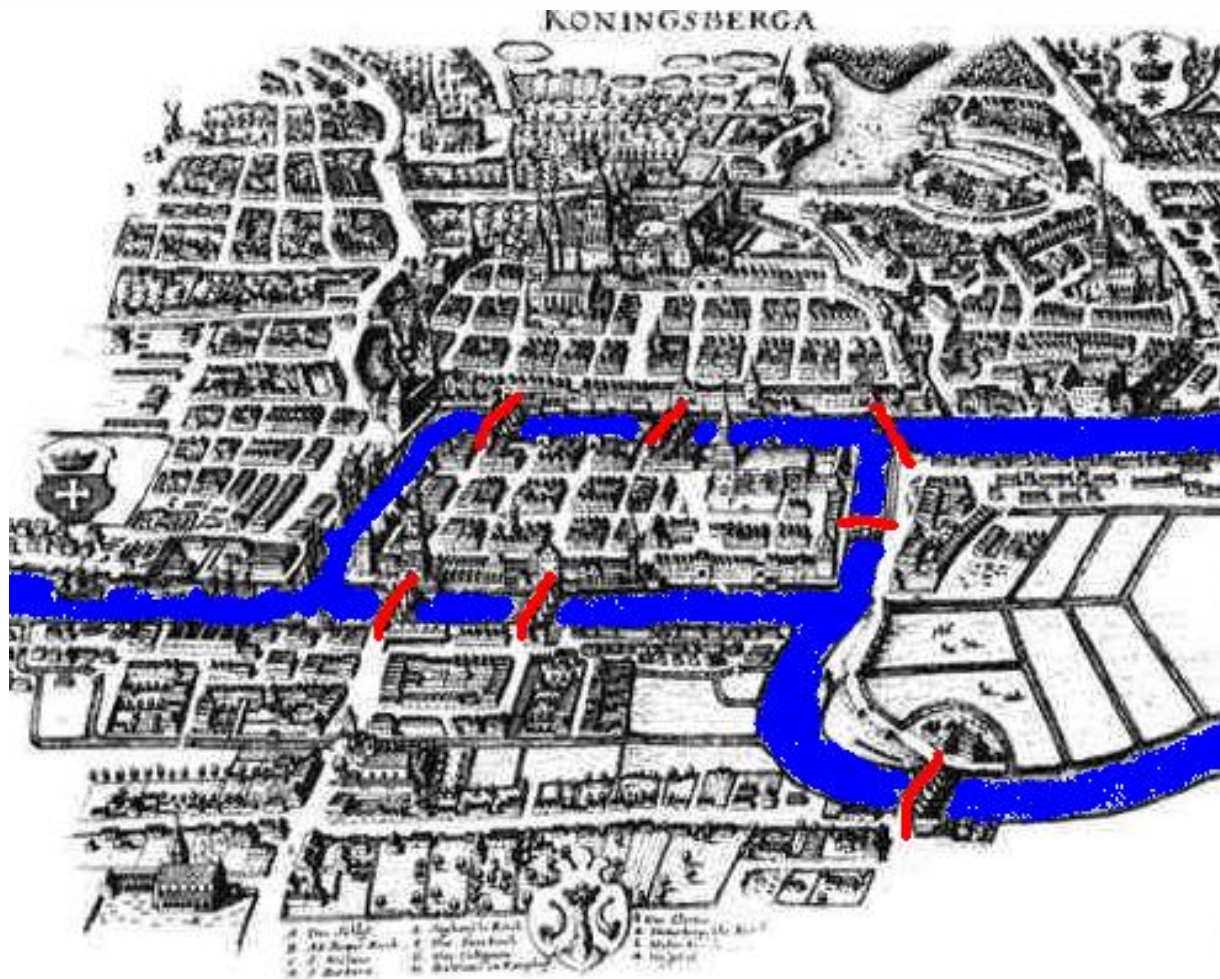
Objectives

- Understand graph terminology
- Implement graphs using
 - Adjacency lists and
 - Adjacency matrices
- Perform graph searches
 - Depth first search
 - Breadth first search
- Perform shortest-path algorithms
 - Disjkstra's algorithm
 - A* algorithm

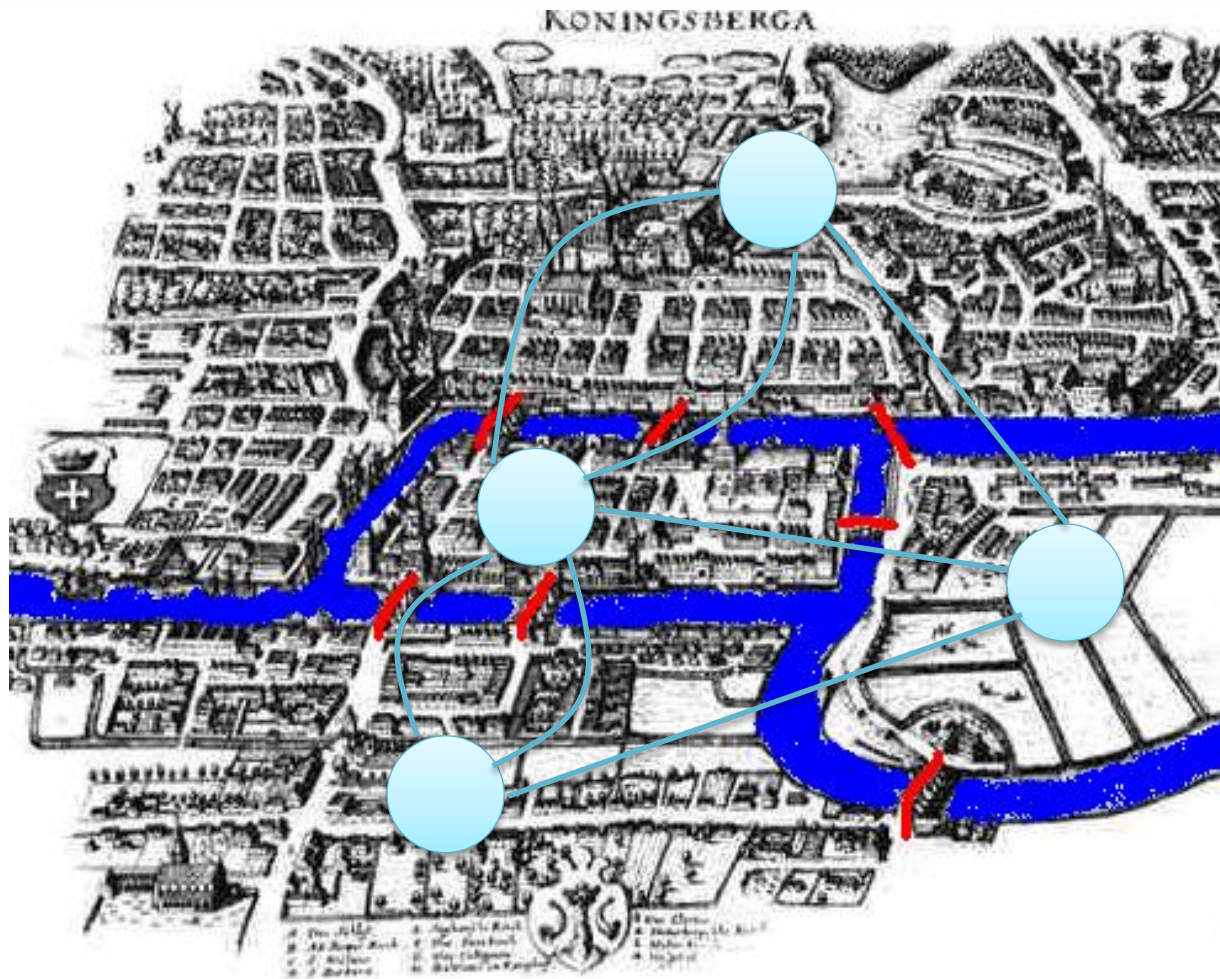
Graph Theory and Euler

- Graph theory is often considered to have been born with Leonhard Euler
 - In 1736 he solved the *Konigsberg bridge problem*
- Konigsberg was a city in Eastern Prussia
 - Renamed Kalinigrad when East Prussia was divided between Poland and Russia in 1945
 - Konigsberg had seven bridges in its centre
 - The inhabitants of Konigsberg liked to see if it was possible to walk across each bridge just once
 - And then return to where they started
 - Euler proved that it was impossible to do this, as part of this proof he represented the problem as a graph

Konigsberg Graph




Konigsberg



Multigraphs

- The Königsberg graph is an example of a *multigraph*
- A multigraph has multiple edges between the same pair of vertices
- In this case the edges represent bridges



Graph Uses

- Graphs are used as representations of many different types of problems
 - Network configuration
 - Airline flight booking
 - Pathfinding algorithms
 - Database dependencies
 - Task scheduling
 - Critical path analysis
 - ...

Graph Terminology


- A graph consists of two sets
 - A set V of *vertices* (or nodes) and
 - A set E of *edges* that connect vertices
 - $|V|$ is the size of V , $|E|$ the size of E
- Two vertices may be connected by a *path*
 - A sequence of edges that begins at one vertex and ends at the other
 - A *simple path* does not pass through the same vertex more than once
 - A *cycle* is a path that starts and ends at the same vertex

Numbers of Vertices and Edges

- If a graph has v vertices, how many edges does it have?
 - If every vertex is connected to every other vertex, and we count each direction as two edges
 - $v^2 - v$
 - If the graph is a tree
 - $v - 1$
 - Minimum number of edges
 - 0


Connected and Unconnected Graphs

- A *connected* graph is one where every pair of distinct vertices has a *path* between them
- A *complete* graph is one where every pair of vertices has an *edge* between them
- A graph cannot have multiple edges between the same pair of vertices
- A graph cannot have *self edges*, an edge from and to the same vertex




connected graph

and a tree




complete graph



unconnected
graph

Directed Graphs


- In a *directed graph* (or digraph) each edge has a direction and is called a directed edge
- A directed edge can only be traveled in one direction
- A pair of vertices in a digraph may have two edges between them, one in each direction



directed graph

Weighted Graphs

- In a *weighted graph* each edge is assigned a weight
 - Edges are labeled with their weights
- Each edge's weight represents the cost to travel along that edge
 - The cost could be distance, time, money or some other measure
 - The cost depends on the underlying problem



weighted graph

Basic Graph Operations

- Create an empty graph
- Test to see if a graph is empty
- Determine the number of vertices in a graph
- Determine the number of edges in a graph
- Determine if an edge exists between two vertices
 - and in a weighted graph determine its weight
- Insert a vertex
 - each vertex is assumed to have a distinct search key
- Delete a vertex, and its associated edges
- Delete an edge
- Return a vertex with a given key


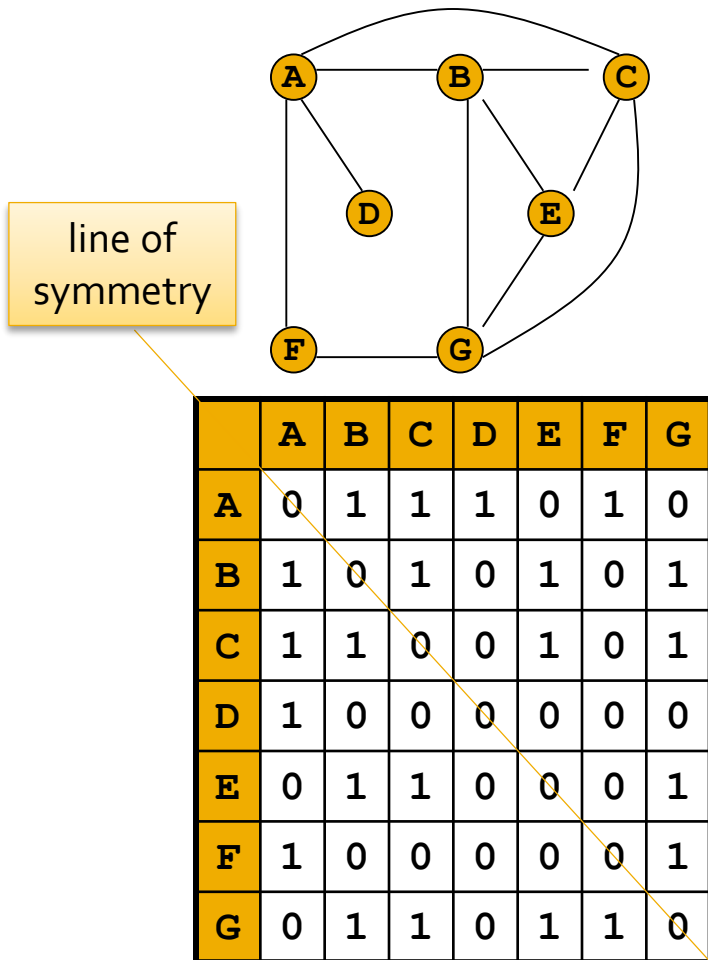
Graph Implementation

- There are two common implementations of graphs
 - Both implementations require a list of all vertices in the set of vertices, V
 - The implementations differ in how edges are recorded
- Adjacency matrices
 - Provide fast lookup of individual edges
 - But waste space for sparse graphs
- Adjacency lists
 - Are more space efficient for sparse graphs
 - Can efficiently find all the neighbours of a vertex

Adjacency Matrix

- The edges are recorded in an $|V| * |V|$ matrix
- In an unweighted graph entries in the matrix are
 - 1 when there is an edge between vertices or
 - 0 when there is no edge between vertices
- In a weighted graph entries are either
 - The edge weight if there is an edge between vertices
 - Infinity when there is no edge between vertices
- Adjacency matrix performance
 - Looking up an edge requires $O(1)$ time
 - Finding all neighbours of a vertex requires $O(|V|)$ time
 - The matrix requires $|V|^2$ space

Adjacency Matrix Examples




	A	B	C	D	E	F	G
A	∞	1	∞	3	∞	5	∞
B	∞	∞	∞	∞	2	∞	∞
C	5	1	∞	∞	∞	∞	∞
D	1	∞	∞	∞	∞	∞	∞
E	∞	∞	2	∞	∞	∞	3
F	∞	∞	∞	∞	∞	∞	8
G	∞	2	4	∞	∞	∞	∞


Adjacency Lists

- The edges are recorded in an array $|V|$ of linked lists
- In an unweighted graph a list at index i records the keys of the vertices adjacent to vertex i
- In a weighted graph a list at index i contains pairs
 - Which record vertex keys (of vertices adjacent to i)
 - And their associated edge weights
- Adjacency List Performance
 - Looking up an edge requires time proportional to the average number of edges
 - Finding all vertices adjacent to a given vertex also takes time proportional to the average number of edges
 - The list requires $O(|E|)$ space

Adjacency List Examples



A	B	→	C	→	D	→	F	
B	A	→	C	→	E	→	G	
C	A	→	B	→	E	→	G	
D	A							
E	B	→	C	→	G			
F	A	→	G					
G	B	→	C	→	E	→	F	




A	B	1	→	D	3	→	F	5	
B	E	2							
C	A	5	→	B	1				
D	A	1							
E	C	2	→	G	3				
F	G	8							
G	B	2	→	C	4				

Graph Traversals


- A graph traversal algorithm visits all of the vertices that can be reached
 - If the graph is not connected some of the vertices will not be visited
 - Therefore a graph traversal algorithm can be used to see if a graph is connected
- Vertices should be marked as *visited*
 - Otherwise, a traversal will go into an infinite loop if the graph contains a cycle

Breadth First Search

- After visiting a vertex, v , visit every vertex adjacent to v before moving on
- Use a queue to store nodes
 - Queues are FIFO
- BFS:
 - visit and insert start
 - while (q not empty)
 - remove node from q and make it *current*
 - visit and insert the unvisited nodes adjacent to *current*



Breadth First Search Example



queue


A
B
F
G
C
H
I
J
D
K
E

visited


A
B
F
G
C
H
I
J
D
K
E

Depth First Search

- Visit a vertex, v , move from v as deeply as possible
- Use a stack to store nodes
 - Stacks are LIFO
- DFS:
 - visit and push start
 - while (s not empty)
 - peek at node, nd , at top of s
 - if nd has an unvisited neighbour visit it and push it onto s
 - else pop nd from s



Depth First Search Example



stack


J K
E I
D H
C G
B F
A

visited

A
B
C
D
E
F
G
H
I
J
K

Shortest Path Problem

- What is the least cost path from one vertex to another?
 - Referred to as the shortest path between vertices
 - For weighted graphs this is the path that has the smallest sum of its edge weights
- Dijkstra's algorithm finds the shortest path between one vertex and all other vertices
 - The algorithm is named after its discoverer, Edgser Dijkstra



The shortest path between B and G is: B-D-E-F-G and not B-G (or B-A-E-F-G)